

document

Load dataset

```
library(readr)
database <- read_rds("../data/database.rds")

colnames(database)
```

```
[1] "Estación"           "Temperatura (°C)"
[3] "Humedad (%)"        "Presión (hPa)"
[5] "Velocidad de viento (m/s)" "CO (mg/m3)"
[7] "NO (ug/m3)"          "NO2 (ug/m3)"
[9] "NOX (ug/m3)"         "O3 (ug/m3)"
[11] "PM10 (ug/m3)"        "time"
```

```
nrow(database)
```

```
[1] 1105
```

```
summary(database)
```

Estación	Temperatura (°C)	Humedad (%)	Presión (hPa)
Length:1105	Min. : 0.00	Min. : 3.00	Min. :902.0
Class :character	1st Qu.:11.40	1st Qu.:32.00	1st Qu.:923.2
Mode :character	Median :18.00	Median :46.00	Median :927.3
	Mean :17.32	Mean :48.04	Mean :926.9
	3rd Qu.:23.50	3rd Qu.:62.00	3rd Qu.:930.5
	Max. :36.60	Max. :95.00	Max. :940.0

Velocidad de viento (m/s)	CO (mg/m3)	NO (ug/m3)	NO2 (ug/m3)
Min. : 0.000	Min. :0.000	Min. : 0.02	Min. : 1.28
1st Qu.: 0.310	1st Qu.:0.740	1st Qu.: 3.53	1st Qu.: 30.71
Median : 1.110	Median :1.120	Median : 26.00	Median : 45.56
Mean : 1.484	Mean :1.172	Mean : 35.36	Mean : 44.17
3rd Qu.: 1.940	3rd Qu.:1.580	3rd Qu.: 55.15	3rd Qu.: 54.11
Max. :18.060	Max. :3.460	Max. :363.27	Max. :131.05

NOX (ug/m3)	O3 (ug/m3)	PM10 (ug/m3)
Min. : 4.15	Min. : 0.00	Min. : 0.00
1st Qu.: 40.11	1st Qu.: 13.05	1st Qu.: 15.00
Median : 93.49	Median : 30.55	Median : 29.00
Mean : 98.40	Mean : 51.02	Mean : 36.75
3rd Qu.:137.60	3rd Qu.: 43.86	3rd Qu.: 49.00
Max. :559.91	Max. :487.52	Max. :318.00
		NA's :11

time
Min. :2022-02-11 17:00:00.00
1st Qu.:2022-05-02 10:00:00.00
Median :2023-01-13 02:00:00.00
Mean :2022-11-29 21:23:43.71
3rd Qu.:2023-03-27 15:00:00.00
Max. :2023-10-17 07:00:00.00

Cleaning

- Tener en cuenta Temperatura, Humedad relativa, Presión atmosférica, Velocidad de viento, CO, NO, NO2, O3 como variables predictoras de PM10.
- Remuevo NA

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(tidyr)
data <- database %>% select(-one_of(c("Estación","time")))
data <- data[complete.cases(data),]
```

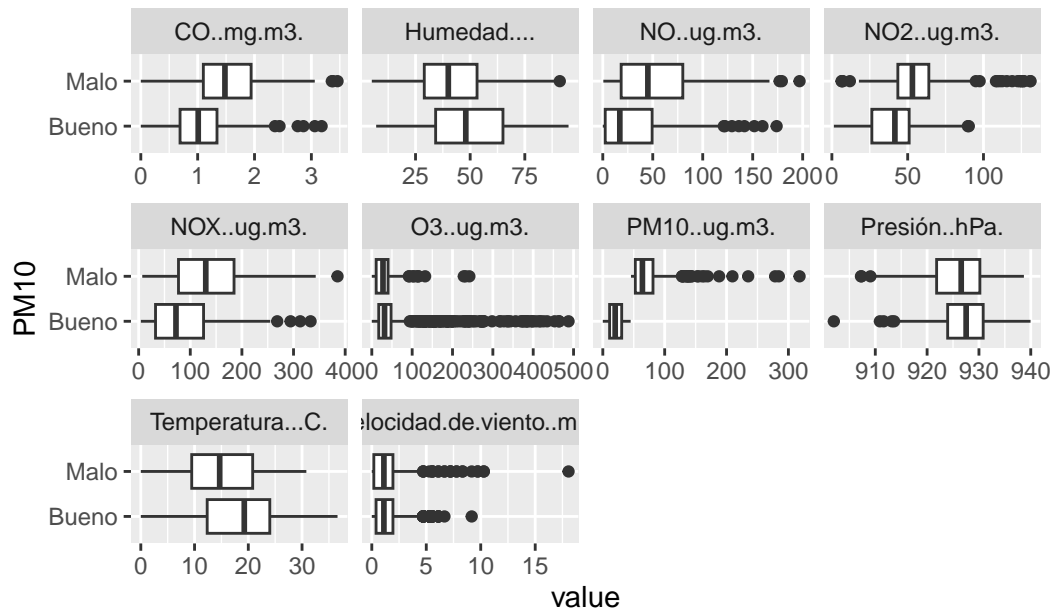
```
nrow(data)
```

```
[1] 1094
```

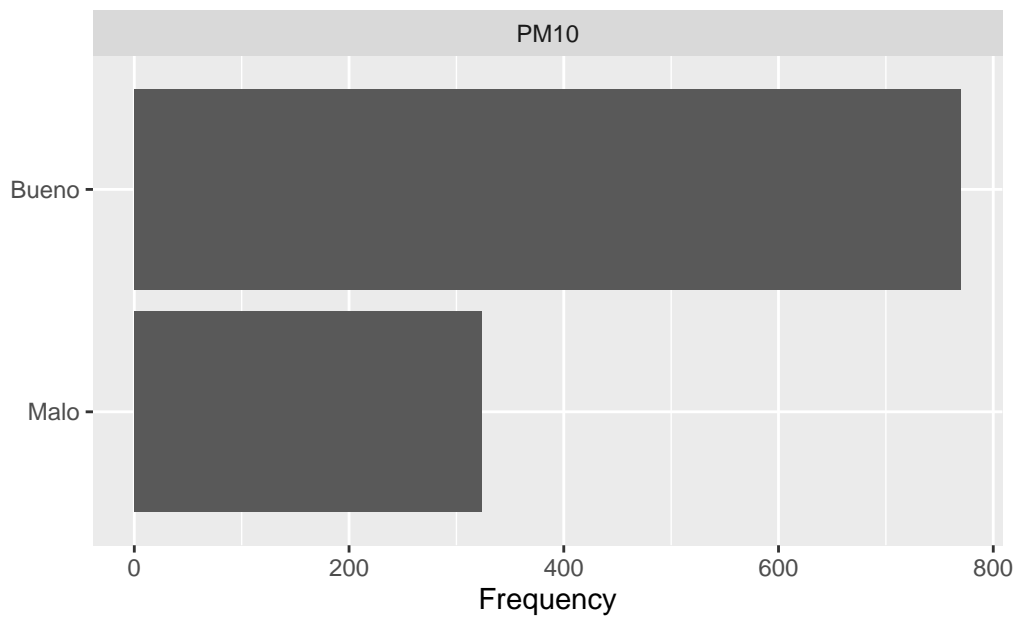
```
summary(data)
```

Temperatura (°C)	Humedad (%)	Presión (hPa)	Velocidad de viento (m/s)
Min. : 0.00	Min. : 5.0	Min. : 902.0	Min. : 0.000
1st Qu.: 11.31	1st Qu.: 32.0	1st Qu.: 923.3	1st Qu.: 0.310
Median : 17.90	Median : 46.0	Median : 927.3	Median : 1.110
Mean : 17.25	Mean : 48.2	Mean : 927.0	Mean : 1.483
3rd Qu.: 23.45	3rd Qu.: 62.0	3rd Qu.: 930.5	3rd Qu.: 1.940
Max. : 36.60	Max. : 95.0	Max. : 940.0	Max. : 18.060
CO (ug/m3)	NO (ug/m3)	NO2 (ug/m3)	NOX (ug/m3)
Min. : 0.000	Min. : 0.020	Min. : 1.28	Min. : 4.15
1st Qu.: 0.740	1st Qu.: 3.542	1st Qu.: 30.94	1st Qu.: 40.31
Median : 1.120	Median : 26.045	Median : 45.73	Median : 94.02
Mean : 1.174	Mean : 35.203	Mean : 44.38	Mean : 98.36
3rd Qu.: 1.577	3rd Qu.: 55.203	3rd Qu.: 54.13	3rd Qu.: 137.90
Max. : 3.460	Max. : 196.860	Max. : 131.05	Max. : 385.19
O3 (ug/m3)	PM10 (ug/m3)		
Min. : 0.00	Min. : 0.00		
1st Qu.: 13.05	1st Qu.: 15.00		
Median : 30.61	Median : 29.00		
Mean : 51.27	Mean : 36.75		
3rd Qu.: 43.81	3rd Qu.: 49.00		
Max. : 487.52	Max. : 318.00		

- Discretizar la variable Material Particulado (PM10) tomando como umbral el valor de 45 µg/m3, por debajo del cual se categorizará como “Bueno”. Por encima de 45 µg/m3, se asignará el valor “Malo”.



```
plot_bar(data)
```



```
data <- data %>% select(-one_of(c("PM10 (ug/m3)")))
```

Modelos

Separación de sets de datos

```
library(tidymodels)
```

```
-- Attaching packages ----- tidymodels 1.2.0 --
```

v broom	1.0.5	v recipes	1.0.10
v dials	1.2.1	v rsample	1.2.1
v ggplot2	3.5.0	v tibble	3.2.1
v infer	1.0.7	v tune	1.2.1
v modeldata	1.3.0	v workflows	1.1.4
v parsnip	1.2.1	v workflowsets	1.1.0
v purrr	1.0.2	v yardstick	1.3.1

```
-- Conflicts ----- tidymodels_conflicts() --
```

```
x purrr::discard() masks scales::discard()
x dplyr::filter() masks stats::filter()
x dplyr::lag() masks stats::lag()
x yardstick::spec() masks readr::spec()
x recipes::step() masks stats::step()
* Learn how to get started at https://www.tidymodels.org/start/
```

```
set.seed(123)
splits <- initial_split(data, strata = PM10, prop = 3/4)

data_train <- training(splits) # 75 % entrenamiento
data_test <- testing(splits) # 25 % en testeo
```

Clasificación binaria

Regresión logística

```
lr_mod <-  
  logistic_reg(penalty = tune(), mixture = 1) %>%  
  set_engine("glmnet")
```

Receta

```
lr_recipe <-  
  recipe(PM10 ~ ., data = data_train) %>%  
  step_normalize(all_predictors())
```

Grid tuning

```
lr_workflow <-  
  workflow() %>%  
  add_model(lr_mod) %>%  
  add_recipe(lr_recipe)
```

Since we have only one hyperparameter to tune here, we can set the grid up manually using a one-column tibble with 30 candidate values:

```
lr_reg_grid <- tibble(penalty = 10^seq(-4, -1, length.out = 30))
```

```
lr_reg_grid
```

```
# A tibble: 30 x 1  
  penalty  
  <dbl>  
1 0.0001  
2 0.000127  
3 0.000161  
4 0.000204  
5 0.000259  
6 0.000329  
7 0.000418  
8 0.000530  
9 0.000672
```

```
10 0.000853
# i 20 more rows
```

Conjunto de validación para usar durante el entrenamiento

```
set.seed(234)
# 20 %
val_set <- validation_split(data_train,
                             strata = PM10,
                             prop = 0.80)
```

Warning: `validation_split()` was deprecated in rsample 1.2.0.
i Please use `initial_validation_split()` instead.

```
lr_res <-
  lr_workflow %>%
  tune_grid(val_set,
            grid = lr_reg_grid,
            control = control_grid(save_pred = TRUE),
            metrics = metric_set(roc_auc))
```


```
lr_res
```

```
# Tuning results
# Validation Set Split (0.8/0.2) using stratification
# A tibble: 1 x 5
  splits          id      .metrics      .notes      .predictions
<list>         <chr>    <list>      <list>      <list>
1 <split [655/165]> validation <tibble [30 x 5]> <tibble [0 x 3]> <tibble>
```

```
lr_plot <-
  lr_res %>%
  collect_metrics() %>%
  ggplot(aes(x = penalty, y = mean)) +
  geom_point() +
  geom_line() +
  ylab("Area under the ROC Curve") +
  scale_x_log10(labels = scales::label_number())
```



```
lr_plot
```



Mejores modelos de Logistic Regression

```
top_models <-  
  lr_res %>%  
  show_best(metric = "roc_auc", n = 15) %>%  
  arrange(penalty)  
top_models
```

A tibble: 15 x 7

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.0001	roc_auc	binary	0.859	1	NA	Preprocessor1_Model01
2	0.000127	roc_auc	binary	0.859	1	NA	Preprocessor1_Model02
3	0.000161	roc_auc	binary	0.859	1	NA	Preprocessor1_Model03
4	0.000204	roc_auc	binary	0.859	1	NA	Preprocessor1_Model04
5	0.000259	roc_auc	binary	0.859	1	NA	Preprocessor1_Model05
6	0.000329	roc_auc	binary	0.859	1	NA	Preprocessor1_Model06
7	0.000418	roc_auc	binary	0.859	1	NA	Preprocessor1_Model07

8	0.000530	roc_auc	binary	0.859	1	NA	Preprocessor1_Model108
9	0.000853	roc_auc	binary	0.859	1	NA	Preprocessor1_Model110
10	0.00108	roc_auc	binary	0.859	1	NA	Preprocessor1_Model111
11	0.00574	roc_auc	binary	0.860	1	NA	Preprocessor1_Model118
12	0.00728	roc_auc	binary	0.861	1	NA	Preprocessor1_Model119
13	0.00924	roc_auc	binary	0.860	1	NA	Preprocessor1_Model120
14	0.0117	roc_auc	binary	0.859	1	NA	Preprocessor1_Model121
15	0.0149	roc_auc	binary	0.860	1	NA	Preprocessor1_Model122

```
lr_best <-
  lr_res %>%
  collect_metrics() %>%
  arrange(penalty) %>%
  slice(12) # modelo 12
```


```
lr_best
```

A tibble: 1 x 7

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.00137	roc_auc	binary	0.859	1	NA	Preprocessor1_Model112

```
lr_auc <-
  lr_res %>%
  collect_predictions(parameters = lr_best) %>%
  roc_curve(PM10, .pred_Bueno) %>%
  mutate(model = "Logistic Regression")

autoplot(lr_auc)
```



Random Forest

```
cores <- parallel::detectCores()
#cores

rf_mod <-
  rand_forest(mtry = tune(), min_n = tune(), trees = 100) %>%
  set_engine("ranger", num.threads = cores) %>%
  set_mode("classification")

rf_recipe <-
  recipe(PM10 ~ ., data = data_train)

rf_workflow <-
  workflow() %>%
  add_model(rf_mod) %>%
  add_recipe(rf_recipe)
```

```

set.seed(345)
rf_res <-
  rf_workflow %>%
    tune_grid(val_set,
              grid = 25,
              control = control_grid(save_pred = TRUE),
              metrics = metric_set(roc_auc))

```

i Creating pre-processing data to finalize unknown parameter: mtry

```

rf_res %>%
  show_best(metric = "roc_auc")

```

```

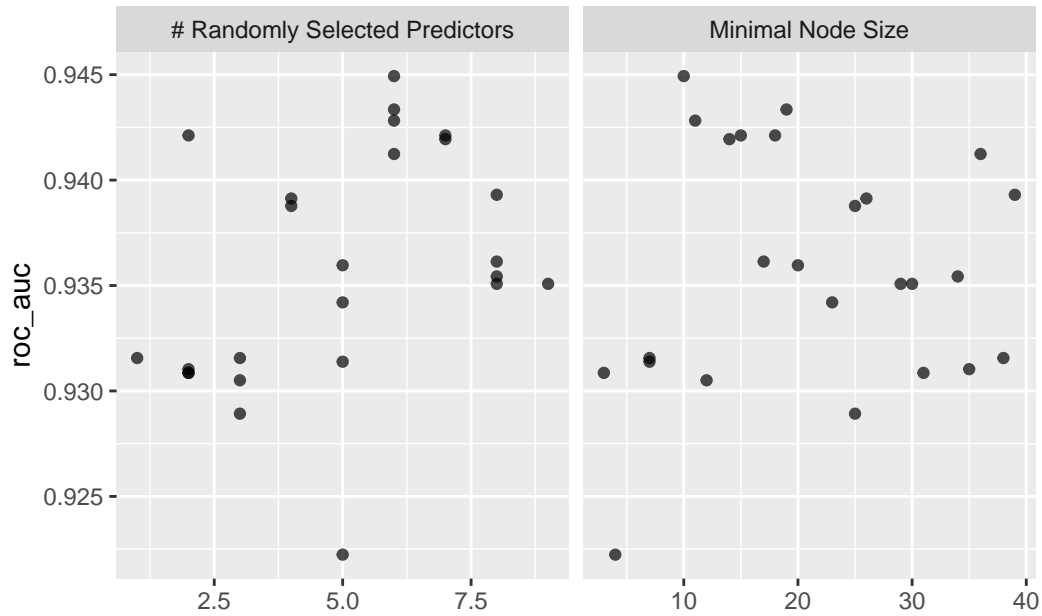
# A tibble: 5 x 8
  mtry min_n .metric .estimator mean      n std_err .config
<int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
1     6    10 roc_auc binary    0.945     1      NA Preprocessor1_Model06
2     6    19 roc_auc binary    0.943     1      NA Preprocessor1_Model10
3     6    11 roc_auc binary    0.943     1      NA Preprocessor1_Model09
4     7    15 roc_auc binary    0.942     1      NA Preprocessor1_Model20
5     2    18 roc_auc binary    0.942     1      NA Preprocessor1_Model24

```

```

autoplot(rf_res)

```



```
rf_best <-
  rf_res %>%
    select_best(metric = "roc_auc")
rf_best
```


```
# A tibble: 1 x 3
  mtry min_n .config
<int> <int> <chr>
1     6    10 Preprocessor1_Model106
```

To filter the predictions for only our best random forest model, we can use the `parameters` argument and pass it our tibble with the best hyperparameter values from tuning, which we called `rf_best`:

```
rf_auc <-
  rf_res %>%
    collect_predictions(parameters = rf_best) %>%
    roc_curve(PM10, .pred_Bueno) %>%
    mutate(model = "Random Forest")

bind_rows(rf_auc, lr_auc) %>%
  ggplot(aes(x = 1 - specificity, y = sensitivity, col = model)) +
```

```
geom_path(lwd = 1.5, alpha = 0.8) +
geom_abline(lty = 3) +
coord_equal() +
scale_color_viridis_d(option = "plasma", end = .6)
```



The random forest is uniformly better across event probability thresholds.

last random forest fit

```
# the last model
last_rf_mod <-
  rand_forest(mtry = 6, min_n = 10, trees = 100) %>%
  set_engine("ranger", num.threads = cores, importance = "impurity") %>%
  set_mode("classification")

# the last workflow
last_rf_workflow <-
  rf_workflow %>%
  update_model(last_rf_mod)

# the last fit
```

```

set.seed(345)
last_rf_fit <-
  last_rf_workflow %>%
  last_fit(splits)

last_rf_fit

```

```

# Resampling results
# Manual resampling
# A tibble: 1 x 6
  splits          id          .metrics .notes   .predictions .workflow
  <list>         <chr>         <list>  <list>   <list>        <list>
1 <split [820/274]> train/test split <tibble> <tibble> <tibble>    <workflow>

```

Vip Variable importance

```
library(vip)
```

Attaching package: 'vip'

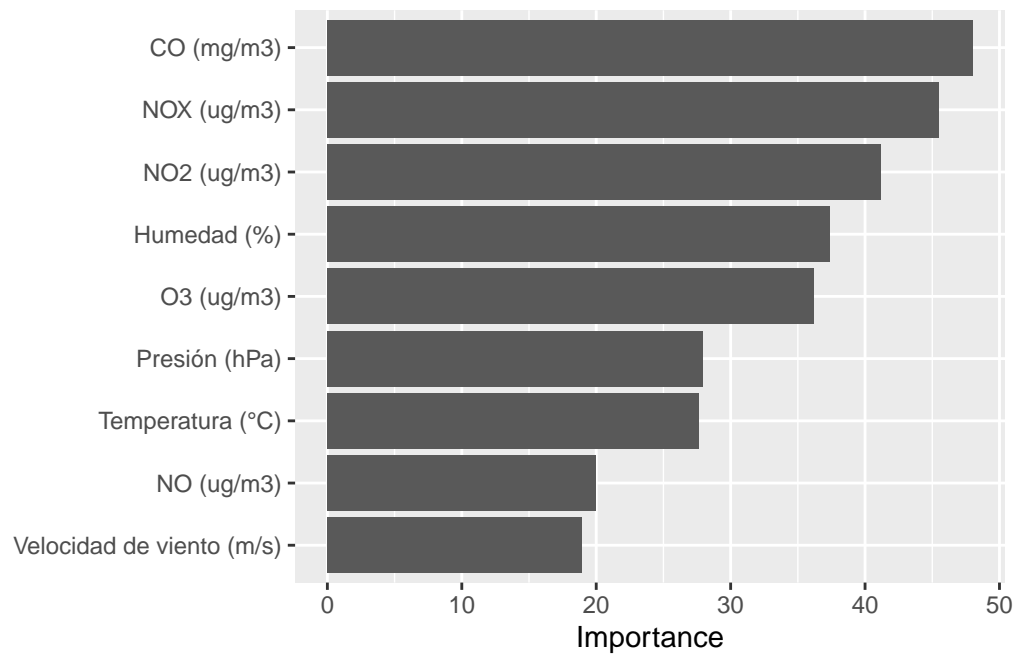
The following object is masked from 'package:utils':

```
vi
```

```

last_rf_fit %>%
  extract_fit_parsnip() %>%
  vip(num_features = 10)

```



```
last_rf_fit %>%  
  collect_predictions() %>%  
  roc_curve(PM10, .pred_Bueno) %>%  
  autoplot()
```